

NoSQL - MongoDB database

Yahyaoui Med Aziz | azizyahyaoui@gmail.com
(<mailto:azizyahyaoui@gmail.com>) | 12122023.

1- Definition

MongoDB is a popular open-source NoSQL (non-relational) database management system. It is designed to handle large amounts of data and provides high performance, scalability, and flexibility. MongoDB stores data in a flexible, JSON-like format called BSON (Binary JSON), making it easy to work with and suitable for a variety of applications.

2- Key features of MongoDB

- **Document-Oriented:** MongoDB stores data in flexible, JSON-like documents called BSON documents. These documents can have nested arrays and subdocuments, allowing for complex data structures.
- **Schema-less:** Unlike traditional relational databases, MongoDB is schema-less. This means that each document in a collection can have a different structure, providing flexibility in data modeling.
- **Query Language:** MongoDB supports a rich query language, allowing for complex queries and filtering of data. It also supports indexing for efficient data retrieval.
- **Scalability:** MongoDB is designed to scale horizontally, allowing you to distribute data across multiple servers or clusters. This makes it suitable for handling large datasets and high traffic loads.
- **Aggregation Framework:** MongoDB includes a powerful aggregation framework that allows for data transformation and analysis within the database itself. It supports a wide range of operations for filtering, grouping, sorting, and projecting data.
- **Replication:** MongoDB supports replication, providing high availability by maintaining multiple copies of data across different servers or clusters. If one server fails, another can take over, ensuring continuous operation.
- **Sharding:** MongoDB supports sharding, which is the horizontal partitioning of data across multiple servers or clusters. This helps distribute the load and enables horizontal scaling for large datasets.
- **Official Drivers:** MongoDB provides official drivers for various programming languages, making it easy for developers to interact with the database using their preferred programming language.

- MongoDB is commonly used in web applications, content management systems, real-time analytics, and other scenarios where a flexible and scalable database solution is required.
-

3- Binary JSON (BSON)

In the context of MongoDB, "BSON" stands for Binary JSON. BSON is a binary representation format used by MongoDB to store and transmit data. It is designed to be efficient in terms of both storage and processing.

BSON is similar to JSON (JavaScript Object Notation), which is a widely used human-readable data interchange format. However, BSON adds additional data types and optimizations to better support the needs of MongoDB.

Here are some key features and characteristics of BSON:

1. **Binary Format:** BSON is a binary representation of JSON-like documents. It allows for efficient storage and transmission of data, especially when compared to plain text-based JSON.
2. **Rich Data Types:** BSON supports a broader range of data types compared to JSON. In addition to the common data types like strings, numbers, booleans, and arrays, BSON includes additional types such as dates, binary data, regular expressions, and more.
3. **Efficient Encoding:** BSON uses a compact binary encoding scheme that allows for efficient serialization and deserialization of data. This makes it suitable for high-performance scenarios, such as database storage and network communication.
4. **Support for Object IDs:** BSON includes a specific data type called "ObjectID" that is used as a unique identifier for documents within a MongoDB collection. ObjectIDs are generated by MongoDB and can be used for indexing and querying purposes.
5. **Nesting and Document Structure:** BSON supports nested documents, allowing for complex and hierarchical data structures. This is particularly useful when storing and querying data in MongoDB's document-oriented model.
6. **Support for Querying and Indexing:** MongoDB's query language, which is based on JSON-like syntax, is designed to work seamlessly with BSON documents. BSON supports indexing, which allows for efficient retrieval and querying of data in MongoDB.

Overall, BSON provides a binary representation of JSON-like documents that is optimized for MongoDB's storage and querying capabilities. It combines the flexibility and ease of use of JSON with the efficiency and additional features required for a database system like MongoDB.

4- MongoDB data types

MongoDB supports various data types, including:

- **String:** Represents a sequence of characters. MongoDB supports UTF-8 encoded strings.

- Integer: Represents a whole number without a fractional component. It can be either 32-bit or 64-bit, depending on the platform.
- Double: Represents a floating-point number with a fractional component. It is stored as 64-bit floating-point.
- Boolean: Represents a boolean value, either true or false.

Date: Represents a date and time value. It stores the date and time in milliseconds since the Unix epoch (January 1, 1970).

- Array: Represents an ordered list of values. Each value in the array can be of any BSON data type, including other arrays.
- Object: Represents a nested document or an embedded document. It consists of key-value pairs, where the keys are strings and the values can be of any BSON data type.
- Null: Represents a null or empty value.
- ObjectId: Represents a unique identifier for a document. It is a 12-byte identifier typically generated by the MongoDB driver.
- Binary: Represents binary data, such as images or files. It can store data in various formats, including generic binary, UUID, MD5, and user-defined types.
- Regular Expression: Represents a pattern used for pattern matching in string fields.
- Decimal128: Represents a decimal floating-point number with higher precision than the Double type. It is stored as a 128-bit decimal floating-point.
- Timestamp: Represents a timestamp value, typically used for internal MongoDB operations.
- MinKey and MaxKey: Represents the lowest and highest possible key values, respectively. They are used for comparisons and sorting.

5- MongoDB Query Language (MQL)

MQL (MongoDB Query Language) is the query language used to interact with MongoDB databases. MQL is designed to query and manipulate documents in MongoDB collections. MongoDB queries are expressed as JSON-like documents, which makes them easy to read and write.

Here are some MongoDB commands for managing users and databases, creating collections, and performing CRUD operations:

1. Managing Users:

- Create a user:

```
db.createUser({
  user: "username",
  pwd: "password",
  roles: ["role1", "role2"]
})
```

- Update a user's password:

```
db.changeUserPassword("username", "newpassword")
```

- Delete a user:

```
db.dropUser("username")
```

2. Managing Databases:

- Create a database:

```
use mydatabase
```

- Show all databases:

```
show dbs
```

- Switch to a different database:

```
use otherdatabase
```

- Delete a database:

```
db.dropDatabase()
```

3. Creating a Collection:

- Create a collection:

```
db.createCollection("mycollection")
```

- Create a collection with options:

```
db.createCollection("mycollection", {
  capped: true,
  size: 1024,
  max: 1000
})
```

4. CRUD Operations:

- Insert documents into a collection:

```
db.mycollection.insertOne({ name: "John", age: 30 })
db.mycollection.insertMany([
  { name: "Alice", age: 25 },
  { name: "Bob", age: 35 }
])
```

- Find documents in a collection:

```
db.mycollection.find({ age: { $gt: 20 } })
```

- Update documents in a collection:

```
db.mycollection.updateOne(  
  { name: "John" },  
  { $set: { age: 31 } }  
)
```

- Delete documents from a collection:

```
db.mycollection.deleteOne({ name: "Alice" })
```

- Delete all documents from a collection:

```
db.mycollection.deleteMany({})
```

- Drop a collection:

```
db.mycollection.drop()
```

These are some of the basic MongoDB commands for managing users and databases, creating collections, and performing CRUD operations.

Here are some common MQL operations:

1. Querying Documents:

To retrieve documents from a collection based on certain criteria, you use the find method. The criteria are specified as a JSON-like document.

```
// Find documents where the "age" field is equal to 25  
db.users.find({ age: 25 })  
  
// Find documents where the "status" field is "active" and "type" is "admin"  
db.users.find({ status: "active", type: "admin" })  
  
// Using comparison operators  
db.products.find({ price: { $gte: 50 } }) // Greater than or equal to
```

To query data from the 'MyAppDB' database in MongoDB, you can use various ways to execute the db.MyAppDB.find() command.

Here are some examples:

- Basic Find:

```
db.MyAppDB.find()
```

This command returns all documents from the 'MyAppDB' database.

- Query with a Filter:

```
db.MyAppDB.find({ field: value })
```

Replace 'field' with the name of the field you want to filter on, and 'value' with the desired value. This command returns documents that match the specified filter criteria.

- Projection:

```
db.MyAppDB.find({}, { field1: 1, field2: 1 })
```

Replace 'field1' and 'field2' with the names of the fields you want to include in the result. This command returns only the specified fields for each document.

- Sorting:

```
db.MyAppDB.find().sort({ field: 1 })
```

Replace 'field' with the name of the field you want to sort by. Use '1' for ascending order and '-1' for descending order. This command returns documents sorted based on the specified field.

- Limiting the Result:

```
db.MyAppDB.find().limit(10)
```

This command limits the number of documents returned to 10. Adjust the number according to your requirements.

- Skipping:

```
db.MyAppDB.find().skip(1)
```

- Combining Multiple Queries:

```
db.MyAppDB.find({ field1: value1, field2: value2 }).sort({ field: 1 }).limit(10)
```

You can combine multiple query options, such as filtering, sorting, and limiting the result, to retrieve specific documents from the 'MyAppDB' database based on your criteria.

These are some common ways to use the `db.MyAppDB.find()` command to query data from the 'MyAppDB' database in MongoDB. Feel free to customize and combine these options to suit your specific needs.

2. Projection:

You can specify which fields to include or exclude in the query result using projection.

```
```\n// Include only the "name" and "email" fields\ndb.users.find({}, { name: 1, email: 1, _id: 0 })\n\n// Exclude the "password" field\ndb.users.find({}, { password: 0 })\n```
```

### 3. Updating Documents:

To update documents in a collection, you use the update or updateOne method.

```
```\n// Update a document (replace the entire document)\ndb.users.update({ username: "john_doe" }, { $set: { status: "inactive" } })\n\n// Update a single document using updateOne\ndb.users.updateOne({ username: "john_doe" }, { $set: { status: "inactive" } })\n```
```

4. Inserting Documents:

To insert documents into a collection, you use the insert or insertOne method.

```
```\n// Insert a new document\ndb.users.insert({ username: "alice", age: 30, status: "active" })\n\n// Insert a single document using insertOne\ndb.users.insertOne({ username: "bob", age: 25, status: "inactive" })\n```
```

### 5. Deleting Documents:

To delete documents from a collection, you use the remove or deleteOne method.

```
```\n// Remove documents where the "status" is "inactive"\ndb.users.remove({ status: "inactive" })\n\n// Remove a single document using deleteOne\ndb.users.deleteOne({ username: "bob" })\n```
```

6. Aggregation Framework:

MongoDB provides a powerful aggregation framework for performing data transformations and analysis.

```
```\n// Aggregation pipeline to calculate the average age of users\ndb.users.aggregate([\n  { $group: { _id: null, avgAge: { $avg: "$age" } } }\n])\n```
```

The syntax is expressive and flexible, allowing you to perform a wide range of operations on your data.  
Queries will depend on your specific use case and the structure of your data.

## MongoDB query operators

This are common MongoDB query operators with examples:

### 1. Comparison Operators:

- `$eq`: Matches values that are equal to a specified value.

```
db.collection.find({ age: { $eq: 30 } })
```

- `$ne`: Matches values that are not equal to a specified value.

```
db.collection.find({ age: { $ne: 25 } })
```

- `$gt`: Matches values that are greater than a specified value.

```
db.collection.find({ quantity: { $gt: 10 } })
```

- `$gte`: Matches values that are greater than or equal to a specified value.

```
db.collection.find({ quantity: { $gte: 20 } })
```

- `$lt`: Matches values that are less than a specified value.

```
db.collection.find({ price: { $lt: 50 } })
```

- `$lte`: Matches values that are less than or equal to a specified value.

```
db.collection.find({ price: { $lte: 100 } })
```

- `$in`: Matches values that exist in a specified array.

```
db.collection.find({ category: { $in: ["Electronics",
"Appliances"] } })
```

- `$nin`: Matches values that do not exist in a specified array.

```
db.collection.find({ category: { $nin: ["Clothing", "Shoes"] } })
```

### 2. Logical Operators:

- `$and`: Joins query clauses with a logical AND.

```
db.collection.find({ $and: [{ age: { $gte: 18 } }, { age: { $lt:
30 } }] })
```

- `$or`: Joins query clauses with a logical OR.



```
db.collection.find({ $or: [{ category: "Electronics" }, {
category: "Appliances" }] })
```

- \$not: Inverts the effect of a query expression.

```
db.collection.find({ quantity: { $not: { $gt: 10 } } })
```

- \$nor: Joins query clauses with a logical NOR.

```
db.collection.find({ $nor: [{ category: "Clothing" }, { category:
"Shoes" }] })
```

### 3. Element Operators:

- \$exists: Matches documents that contain a specific field.

```
db.collection.find({ price: { $exists: true } })
```

- \$type: Matches documents based on the BSON data type of a field.

```
db.collection.find({ age: { $type: "number" } })
```

### 4. Array Operators:

- \$all: Matches arrays that contain all the specified elements.

```
db.collection.find({ tags: { $all: ["mongodb", "database"] } })
```

- \$elemMatch: Matches arrays that contain at least one element matching all the specified criteria.

```
db.collection.find({ scores: { $elemMatch: { $gte: 80, $lt: 90 } }
})
```

- \$size: Matches arrays that have a specific number of elements.

```
db.collection.find({ days: { $size: 7 } })
```

### 5. Evaluation Operators:

- \$expr: Allows the use of aggregation expressions within the query language.

```
db.collection.find({ $expr: { $gt: ["$price", "$discountedPrice"]
} })
```

- \$regex: Matches documents based on a specified regular expression pattern.

```
db.collection.find({ name: { $regex: "^A" } })
```

- \$text: Performs text search on text indexes.

```
db.collection.find({ $text: { $search: "mongodb" } })
```

- \$mod: Performs a modulo operation on the field value.

```
db.collection.find({ quantity: { $mod: [5, 0] } })
```

## 6. Geospatial Operators:

- `$geoWithin`: Matches documents that are within a specified geometry.

```
db.collection.find({ location: { $geoWithin: { $centerSphere:
[[-73.9667, 40.78], 5 / 3963.2] } } })
```

- `$geoIntersects`: Matches documents that intersect with a specified geometry.

```
db.collection.find({ location: { $geoIntersects: { $geometry: {
type: "Point", coordinates: [-73.9667, 40.78] } } })
```

Apologies for the incomplete response. Here are the remaining examples:

- `$near`: Returns documents based on proximity to a specified point.

```
db.collection.find({ location: { $near: { $geometry: { type:
"Point", coordinates: [-73.9667, 40.78] } }, $maxDistance: 1000 } } })
```

## 7. Projection Operators:

- `$elemMatch`: Projects only the first array element that matches the specified criteria.

```
db.collection.find({ scores: { $elemMatch: { $gte: 80, $lt: 90 } } }, { scores: { $elemMatch: { $gte: 80, $lt: 90 } } })
```

- `$slice`: Limits the number of array elements that are returned in the projection.

```
db.collection.find({}, { scores: { $slice: 3 } })
```

## 8. Array Update Operators:

- `$push`: Appends an element to an array field.

```
db.collection.updateOne({ _id: 1 }, { $push: { scores: 90 } })
```

- `$pull`: Removes all array elements that match a specified condition.

```
db.collection.updateOne({ _id: 1 }, { $pull: { scores: { $lt: 70 } } })
```

- `$addToSet`: Adds an element to an array field only if it does not already exist.

```
db.collection.updateOne({ _id: 1 }, { $addToSet: { tags: "new" } })
```

These examples demonstrate how to use the MongoDB query operators in various scenarios to retrieve, filter, and manipulate data in your collections.

---

## 6- Roles and Privileges

In MongoDB, roles are used to define the set of privileges and permissions granted to a user or a group of users. Roles control what actions users can perform on databases, collections, and documents within a MongoDB deployment. Here are some key points about roles in MongoDB:

1. **Built-in Roles:** MongoDB provides several built-in roles that cover common use cases. Some of the commonly used built-in roles include:
  - **read:** Provides read-only access to specific databases or collections.
  - **readWrite:** Provides read and write access to specific databases or collections.
  - **dbAdmin:** Provides administrative privileges for a specific database.
  - **userAdmin:** Provides administrative privileges for managing users and roles within a specific database.
  - **clusterAdmin:** Provides full administrative privileges across the entire MongoDB deployment.
2. **Privileges:** Roles are associated with a set of privileges that define the actions a user can perform. Privileges include read, write, find, insert, update, delete, create, drop, and more. Privileges can be granted at the database level or the collection level.
3. **Role Hierarchy:** Roles in MongoDB can be hierarchical, meaning that a role can inherit privileges from another role. This allows for the creation of more granular roles by building upon existing roles.
4. **Custom Roles:** In addition to the built-in roles, MongoDB allows you to create custom roles tailored to your specific requirements. Custom roles can be defined with specific privileges and can be assigned to users or groups of users.
5. **Role-Based Access Control (RBAC):** MongoDB's role-based access control system allows administrators to manage user access and permissions at a fine-grained level. By assigning appropriate roles to users, you can control their actions on databases, collections, and documents.
6. **Role Assignment:** Roles can be assigned to users at the database level or the cluster level. Database-level roles apply to a specific database, while cluster-level roles apply to the entire MongoDB deployment.
7. **Role Management:** You can create, modify, and delete roles using MongoDB's administrative commands or through MongoDB's administrative tools, such as the MongoDB shell, MongoDB Compass, or MongoDB Atlas.

Roles play a crucial role in securing your MongoDB deployment by ensuring that users have appropriate access and permissions. By defining roles with the necessary privileges, you can control and restrict user actions within your MongoDB databases and collections.

## Grant and Revoke roles and privileges

These are examples of how to grant and revoke roles and privileges using MQL (MongoDB Query Language) commands in Mongosh:

## 1. Granting Roles:

- Grant a built-in role to a user:

```
db.grantRolesToUser("username", ["read"])
```

- Grant multiple built-in roles to a user:

```
db.grantRolesToUser("username", ["read", "write"])
```

- Grant a custom role to a user:

```
db.grantRolesToUser("username", [{ role: "customRole", db: "mydatabase" }])
```

## 2. Revoking Roles:

- Revoke a role from a user:

```
db.revokeRolesFromUser("username", ["read"])
```

- Revoke multiple roles from a user:

```
db.revokeRolesFromUser("username", ["read", "write"])
```

- Revoke a custom role from a user:

```
db.revokeRolesFromUser("username", [{ role: "customRole", db: "mydatabase" }])
```

## 3. Granting Privileges:

- Grant privileges to a user on a specific database:

```
db.grantPrivilegesToUser("username", [
 { resource: { db: "mydatabase", collection: "" }, actions: ["find"] }
])
```

- Grant privileges to a user on a specific collection:

```
db.grantPrivilegesToUser("username", [
 { resource: { db: "mydatabase", collection: "mycollection" }, actions: ["insert", "update"] }
])
```

## 4. Revoking Privileges:

- Revoke privileges from a user on a specific database:

```
db.revokePrivilegesFromUser("username", [
 { resource: { db: "mydatabase", collection: "" }, actions: ["find"] }
])
```

- Revoke privileges from a user on a specific collection:

```
db.revokePrivilegesFromUser("username", [
 { resource: { db: "mydatabase", collection: "mycollection" },
 actions: ["insert", "update"] }
])
```

### Example scenario:

If I've a user called 'OLiver' and a DB named "MyAppDB" how to grant him only to use this DB.

To grant the user 'Oliver' access to only the 'MyAppDB' database in MongoDB, you can use the following steps:

1. Switch to the 'MyAppDB' database:

```
use MyAppDB
```
```

2. Create a user 'Oliver' and grant him the necessary privileges for the 'MyAppDB' database:

```
db.createUser({
  user: 'Oliver',
  pwd: 'password', // Replace with the desired password
  roles: [{ role: 'readWrite', db: 'MyAppDB' }]
})
```
```

This will create a user named 'Oliver' with a password of your choice and grant the 'readWrite' role specifically for the 'MyAppDB' database. The 'readWrite' role allows the user to perform both read and write operations on the database.

Once you've executed these commands, the user 'Oliver' will have access only to the 'MyAppDB' database and will not be able to access any other databases in the MongoDB deployment.

How to revoke him from changing his password

In MongoDB, users with appropriate privileges have the ability to change their own password by default. If you want to revoke the user 'Oliver' from changing his own password, you can follow these steps:

1. Switch to the 'admin' database:

```
use MyAppDB
```
```

2. Revoke the 'changeOwnPassword' privilege from the 'Oliver' user:

```
db.revokePrivilegesFromRole('readWrite', [{ resource: { db: 'admin' },
actions: ['changeOwnPassword'] }])
```
```

This command revokes the 'changeOwnPassword' privilege from the 'readWrite' role associated with the 'Oliver' user. By removing this privilege, the user will no longer be able to change their own password.

Please note that you should have appropriate administrative privileges to perform this action. Additionally, revoking the 'changeOwnPassword' privilege will affect all users with the 'readWrite' role. If you have other users with the 'readWrite' role who should retain the ability to change their own password, you may need to create a separate role specifically for 'Oliver' without the 'changeOwnPassword' privilege.

---

## 7- What are indexes in MongoDB and its job?

In MongoDB, an index is a data structure that improves the speed of data retrieval operations on a database. Indexes store a small subset of the data from a collection and provide an efficient way to locate documents within that collection based on the values of specific fields. MongoDB uses a B-tree data structure for its indexes.

Key Points about Indexes in MongoDB:

- **Improving Query Performance:**  
The primary purpose of indexes is to improve the speed of queries. By creating indexes on fields frequently used in queries, MongoDB can quickly locate the documents that match the query criteria without performing a collection scan.
- **Structure of an Index:**  
Each index is a separate data structure that includes a sorted or hashed list of the indexed field's values and a reference to the corresponding documents in the collection.
- **Automatic Indexing:**  
MongoDB automatically creates an index on the `_id` field, which serves as the primary key for the collection. Additionally, you can manually create indexes on one or more fields of your choice.
- **Single-field and Compound Indexes:**  
MongoDB supports both single-field indexes and compound indexes. Single-field indexes are created on a single field, while compound indexes are created on multiple fields. Compound indexes can be useful when filtering on multiple criteria.
- **Unique Indexes:**  
You can create unique indexes to ensure that a particular field or combination of fields has unique values across all documents in the collection. This is similar to the concept of unique constraints in relational databases.
- **Sparse Indexes:**  
Sparse indexes only include documents that contain the indexed field. This is useful when the indexed field is not present in all documents.

- **Geospatial Indexes:**  
MongoDB supports geospatial indexes for queries involving location-based data. Geospatial indexes enable efficient retrieval of documents based on their proximity to a specified point or within a specified area.
- **Text Indexes:**  
Text indexes are designed for full-text search queries. They allow you to perform text searches on string content within documents.

How Indexes Work:

- **Query Optimization:**

When MongoDB processes a query, it checks if the query can be satisfied using an index. If an appropriate index exists, MongoDB uses it to locate the documents more efficiently.

- **Index Selection:**

The query planner evaluates multiple candidate indexes and selects the most efficient one based on factors such as query complexity, index size, and the number of documents to scan.

- **Covered Queries:**

In some cases, the index itself may contain all the information needed to satisfy a query. These are called covered queries, and they can be more efficient because MongoDB can fulfill the query requirements directly from the index without accessing the actual documents.

Indexes in MongoDB play a crucial role in improving query performance by allowing the database to locate and retrieve relevant documents more efficiently. Choosing the right indexes based on your application's query patterns is an important aspect of MongoDB database design.

## Create an indexes:

In MongoDB, using the `createIndex()` method. You can create indexes on one or more fields, and you have the flexibility to specify options such as whether the index is unique, sparse, or the type of index (ascending or descending). Here are examples of how to create indexes in MongoDB:

1. **Creating a Single-Field Index:**

To create an index on a single field, you can use the `createIndex()` method:

```
// Create a single-field index on the "username" field
db.collection.createIndex({ username: 1 })
```

In this example:

1 indicates that the index should be in ascending order. You can use -1 for descending order. 2. **Creating a Compound Index:**

To create an index on multiple fields, you can pass an object with multiple field names to the `createIndex()` method:

```
// Create a compound index on the "lastName" and "firstName" fields
db.collection.createIndex({ lastName: 1, firstName: 1 }) 3.
```

Creating a Unique Index:

To create a unique index (ensuring that the indexed fields have unique values across all documents), you can use the unique option:

```
// Create a unique index on the "email" field
db.collection.createIndex({ email: 1 }, { unique: true }) 4.
```

Creating a Sparse Index:

To create a sparse index (index only documents that contain the indexed field), you can use the sparse option:

```
// Create a sparse index on the "phone" field
db.collection.createIndex({ phone: 1 }, { sparse: true }) 5.
```

Creating a Text Index:

To create a text index for full-text search queries, you can use the text index type:

```
// Create a text index on the "description" field
db.collection.createIndex({ description: "text" }) 6.
```

Creating a Geospatial Index:

For geospatial queries, you can create a geospatial index using the 2dsphere or 2d index type:

```
// Create a 2dsphere index on the "location" field for geospatial queries
db.collection.createIndex({ location: "2dsphere" })
```

Important Notes:

Index creation can be a resource-intensive operation, so it's typically done during periods of lower activity.

Indexes should be created based on the queries your application performs. Analyze your query patterns to determine which fields should be indexed.

MongoDB automatically creates an index on the `_id` field by

---

## 8- MongoDB schema and relations

In MongoDB, unlike traditional relational databases, there is no strict requirement for a fixed schema. MongoDB uses a flexible, schema-less document model, which allows for dynamic and nested data structures. This flexibility is particularly useful for applications with evolving requirements. MongoDB's data modeling relies on the BSON (Binary JSON) format to represent documents, and it supports various data relationships, including one-to-one, one-to-many, and many-to-many.

### Data Modeling Concepts:

#### 1. Document:

- In MongoDB, data is stored in BSON documents. Each document is a JSON-like structure with key-value pairs, and documents in a collection can have different fields and structures.



## 2. Collection:

- A collection is a grouping of MongoDB documents. It is similar to a table in a relational database. Collections do not enforce a schema, meaning documents within a collection can have different fields.

## Relationships:

### 1. One-to-One Relationship:

- In a one-to-one relationship, a document in one collection is associated with at most one document in another collection. This relationship is often represented by embedding documents or using references.

- **Embedding:**

```
// User document with an embedded profile
{
 "_id": ObjectId("user_id"),
 "username": "john_doe",
 "email": "john@example.com",
 "profile": {
 "firstName": "John",
 "lastName": "Doe",
 "age": 30
 }
}
```

- **Reference:**

```
// User document with a reference to a Profile document
// User document
{
 "_id": ObjectId("user_id"),
 "username": "john_doe",
 "email": "john@example.com",
 "profileId": ObjectId("profile_id")
}

// Profile document
{
 "_id": ObjectId("profile_id"),
 "firstName": "John",
 "lastName": "Doe",
 "age": 30
}
```

### 2. One-to-Many Relationship:

- In a one-to-many relationship, a document in one collection is associated with multiple documents in another collection.
- **Embedding Arrays:**

```
// Post document with embedded comments
{
 "_id": ObjectId("post_id"),
 "title": "My MongoDB Journey",
 "content": "...",
 "comments": [
 { "userId": ObjectId("user1_id"), "text": "Great post!" },
 { "userId": ObjectId("user2_id"), "text": "Informative." }
]
}
```

- **References:**

```
// Post document with references to Comment documents
// Post document
{
 "_id": ObjectId("post_id"),
 "title": "My MongoDB Journey",
 "content": "...",
 "comments": [
 ObjectId("comment1_id"),
 ObjectId("comment2_id")
]
}

// Comment documents
{
 "_id": ObjectId("comment1_id"),
 "userId": ObjectId("user1_id"),
 "text": "Great post!"
}

{
 "_id": ObjectId("comment2_id"),
 "userId": ObjectId("user2_id"),
 "text": "Informative."
}
```

### 3. Many-to-Many Relationship:

- In a many-to-many relationship, documents in one collection are associated with multiple documents in another collection, and vice versa.
- **Embedding Arrays or References:**

```
// Student document with embedded courses
{
 "_id": ObjectId("student_id"),
 "name": "Alice",
 "courses": [
 { "courseId": ObjectId("course1_id"), "name": "Math" },
 { "courseId": ObjectId("course2_id"), "name": "History" }
]
}
```

```
// Course document with references to Student documents
// Course document
{
 "_id": ObjectId("course1_id"),
 "name": "Math",
 "students": [
 ObjectId("student1_id"),
 ObjectId("student2_id")
]
}

// Student documents
{
 "_id": ObjectId("student1_id"),
 "name": "Alice"
}

{
 "_id": ObjectId("student2_id"),
 "name": "Bob"
}
```

### Considerations:

- **Embedding vs. Referencing:**

- Embedding is suitable when the embedded documents are relatively small and frequently accessed together. Referencing is useful for larger documents or when documents are frequently updated independently.

- **Data Access Patterns:**

- Data modeling should be based on the application's data access patterns. Consider how data will be read and written.

- **Atomicity:**

- Atomic operations are performed on a single document. If you need atomic updates across multiple documents, referencing may be more suitable.

- **Query Performance:**

- Indexing is crucial for good query performance. Analyze your queries and create appropriate indexes.

- **Denormalization:**

- Denormalization, or duplicating data, can be used to optimize read performance. However, it comes at the cost of increased storage and potential data inconsistency.

In MongoDB, the choice of schema and relationship representation depends on the specific requirements of your application and the access patterns of your data. Understanding the application's needs and considering factors such as read and write patterns, data size, and performance requirements will guide your decisions in data modeling.

---

## 9- Aggregations and how create a pipeline

MongoDB provides a powerful aggregation framework for data processing and analysis. It allows you to perform complex operations on data within collections. Here are some common aggregation stages and examples:

1. **\$match:** Filters documents based on specific criteria.

```
db.collection.aggregate([
 { $match: { age: { $gte: 18 } } }
])
```
```

2. **\$group:** Groups documents by a specified key and performs calculations on grouped data.

```
db.collection.aggregate([
  { $group: { _id: "$category", total: { $sum: "$quantity" } } }
])
```
```

3. **\$project:** Reshapes the documents, including specific fields or calculated expressions.

```
db.collection.aggregate([
 { $project: { name: 1, price: 1, discountedPrice: { $multiply:
["$price", 0.9] } } }
])
```
```

4. **\$sort:** Sorts the documents based on specified fields.

```
db.collection.aggregate([
  { $sort: { age: 1 } }
])
```
```

5. **\$limit:** Limits the number of documents in the output.

```
db.collection.aggregate([
 { $limit: 10 }
])
```
```

6. \$skip: Skips a specified number of documents in the input.

```
db.collection.aggregate([
  { $skip: 5 }
])
```
```

7. \$unwind: Deconstructs an array field into multiple documents.

```
db.collection.aggregate([
 { $unwind: "$tags" }
])
```
```

8. \$lookup: Performs a left outer join with another collection.

```
db.collection.aggregate([
  {
    $lookup: {
      from: "orders",
      localField: "_id",
      foreignField: "customerId",
      as: "customerOrders"
    }
  }
])
```
```

9. \$group with \$push: Groups documents and creates an array of values for a specified field.

```
db.collection.aggregate([
 {
 $group: {
 _id: "$category",
 products: { $push: "$name" }
 }
 }
])
```
```

10. \$facet: Enables multiple pipelines to run within a single aggregation stage.

```

db.collection.aggregate([
  {
    $facet: {
      categoryCounts: [{ $group: { _id: "$category", count: { $sum: 1 } } }],
      averagePrice: [{ $group: { _id: null, avgPrice: { $avg: "$price" } } }]
    }
  }
])

```

To create an aggregation pipeline, you can chain these stages together to form a sequence of operations. Each stage takes the output of the previous stage as its input. For example:

```

db.collection.aggregate([
  { $match: { age: { $gte: 18 } } },
  { $group: { _id: "$category", total: { $sum: "$quantity" } } },
  { $sort: { total: -1 } },
  { $limit: 5 }
])

```

In this example, the pipeline filters documents with an age greater than or equal to 18, groups the remaining documents by category, calculates the total quantity for each category, sorts the results in descending order based on the total, and finally limits the output to the top 5 categories.

You can customize the pipeline stages according to your specific data processing requirements.

Here are some additional common aggregation stages in MongoDB:

1. `$addField`: Adds new fields to the documents with specified values or expressions.

```

db.collection.aggregate([
  { $addField: { totalPrice: { $multiply: ["$price", "$quantity"] } } }
])
...

```

2. `$project` with `$cond`: Conditionally includes/excludes fields based on specified conditions.

```

db.collection.aggregate([
  {
    $project: {
      name: 1,
      price: 1,
      discountPercentage: {
        $cond: {
          if: { $gte: ["$quantity", 10] },
          then: 10,
          else: 5
        }
      }
    }
  }
])
```

```

3. \$lookup with pipeline: Performs a more complex join operation using a pipeline as the join condition.

```

db.collection.aggregate([
 {
 $lookup: {
 from: "orders",
 let: { productId: "$_id" },
 pipeline: [
 { $match: { $expr: { $eq: ["$productId", "$$productId"] } } },
 { $sort: { date: -1 } },
 { $limit: 1 }
],
 as: "latestOrder"
 }
 }
])
```

```

4. \$group with \$avg: Calculates the average value of a specified field within each group.

```

db.collection.aggregate([
  {
    $group: {
      _id: "$category",
      averagePrice: { $avg: "$price" }
    }
  }
])
```

```

5. \$facet with \$count: Counts the number of documents matching specified conditions within a facet.

```

db.collection.aggregate([
 {
 $facet: {
 categoryCounts: [
 { $match: { quantity: { $gt: 0 } } },
 { $count: "count" }
]
 }
 }
])
```

```

6. `$redact`: Conditionally includes or excludes documents based on specified conditions.

```

db.collection.aggregate([
  {
    $redact: {
      $cond: {
        if: { $eq: ["$status", "active"] },
        then: "$$DESCEND",
        else: "$$PRUNE"
      }
    }
  }
])
```

```

7. `$graphLookup`: Performs recursive graph traversal operations on a collection.

```

db.collection.aggregate([
 {
 $graphLookup: {
 from: "collection",
 startWith: "$_id",
 connectFromField: "parent",
 connectToField: "_id",
 as: "ancestors"
 }
 }
])
```

```

These additional aggregation stages provide more flexibility and power to perform advanced data transformations and calculations within MongoDB. The aggregation framework is highly versatile and can be combined and customized to suit various data analysis scenarios.

10- Connect MongoDB with some programming languages

JavaScript

To use MongoDB with JavaScript, you can utilize the official MongoDB Node.js driver, which provides a set of APIs for interacting with MongoDB from JavaScript. Here's a step-by-step guide on how to use MongoDB with JavaScript:

1. Install the MongoDB Node.js driver:

```
npm install mongodb
```
```

2. Require the MongoDB driver in your JavaScript file:

```
const { MongoClient } = require('mongodb');
```
```

3. Connect to the MongoDB server:

```
const uri = 'mongodb://localhost:27017'; // Replace with your MongoDB
connection string
const client = new MongoClient(uri, { useNewUrlParser: true,
useUnifiedTopology: true });

async function connect() {
  try {
    await client.connect();
    console.log('Connected to MongoDB');
  } catch (error) {
    console.error('Error connecting to MongoDB:', error);
  }
}

connect();
```
```

4. Perform database operations:

Once connected, you can perform various operations like inserting, updating, deleting, and querying documents in MongoDB.

Here's an example of inserting a document into a collection:

```

async function insertDocument(databaseName, collectionName, document) {
 const db = client.db(databaseName);
 const collection = db.collection(collectionName);

 try {
 const result = await collection.insertOne(document);
 console.log('Inserted document:', result.insertedId);
 } catch (error) {
 console.error('Error inserting document:', error);
 }
}

// Usage example
const document = { name: 'John Doe', age: 25 };
insertDocument('mydatabase', 'mycollection', document);
```

```

You can explore other MongoDB operations like updating documents, deleting documents, and querying documents using the MongoDB Node.js driver's API. Refer to the official MongoDB Node.js driver documentation for more details: <https://mongodb.github.io/node-mongodb-native/>

5. Disconnect from the MongoDB server:

```

async function disconnect() {
  try {
    await client.close();
    console.log('Disconnected from MongoDB');
  } catch (error) {
    console.error('Error disconnecting from MongoDB:', error);
  }
}

disconnect();
```

```

Remember to replace 'mongodb://localhost:27017' with your actual MongoDB connection string, which includes the necessary credentials, database name, and host information.

Here are some examples of CRUD operations (Create, Read, Update, Delete) using the MongoDB Node.js driver for a "person" collection with fields like last\_name, first\_name, email, age, DateOfBirth, address (containing street and zip\_code), and profession:

##### 1. Create (Insert) a Person:

```

async function createPerson(person) {
 const db = client.db('mydatabase');
 const collection = db.collection('person');

 try {
 const result = await collection.insertOne(person);
 console.log('Person created:', result.insertedId);
 } catch (error) {
 console.error('Error creating person:', error);
 }
}

// Usage example
const person = {
 last_name: 'Doe',
 first_name: 'John',
 email: 'john.doe@example.com',
 age: 25,
 DateOfBirth: new Date('1998-01-01'),
 address: {
 street: '123 Main St',
 zip_code: '12345'
 },
 profession: 'Engineer'
};

createPerson(person);

```

## 2. Read (Find) a Person:

```

async function findPersonByEmail(email) {
 const db = client.db('mydatabase');
 const collection = db.collection('person');

 try {
 const person = await collection.findOne({ email });
 console.log('Person found:', person);
 } catch (error) {
 console.error('Error finding person:', error);
 }
}

// Usage example
const email = 'john.doe@example.com';
findPersonByEmail(email);

```

## 3. Update a Person:

```

async function updatePersonByEmail(email, updates) {
 const db = client.db('mydatabase');
 const collection = db.collection('person');

 try {
 const result = await collection.updateOne({ email }, { $set: updates });
 console.log('Person updated:', result.modifiedCount);
 } catch (error) {
 console.error('Error updating person:', error);
 }
}

// Usage example
const email = 'john.doe@example.com';
const updates = { age: 26, profession: 'Senior Engineer' };
updatePersonByEmail(email, updates);

```

#### 4. Delete a Person:

```

async function deletePersonByEmail(email) {
 const db = client.db('mydatabase');
 const collection = db.collection('person');

 try {
 const result = await collection.deleteOne({ email });
 console.log('Person deleted:', result.deletedCount);
 } catch (error) {
 console.error('Error deleting person:', error);
 }
}

// Usage example
const email = 'john.doe@example.com';
deletePersonByEmail(email);

```

## Java

To use MongoDB with a Java application, you can utilize the official MongoDB Java driver, which provides a set of APIs for interacting with MongoDB from Java. Here's a step-by-step guide on how to use MongoDB with a Java application:

1. Add the MongoDB Java driver dependency to your project:  
If you're using Maven, add the following dependency to your pom.xml file:

```
<dependency>
 <groupId>org.mongodb</groupId>
 <artifactId>mongodb-driver-sync</artifactId>
 <version>4.4.5</version>
</dependency>
```
```

If you're using Gradle, add the following dependency to your `build.gradle` file:

```
```groovy
implementation 'org.mongodb:mongodb-driver-sync:4.4.5'
```
```

Make sure to check the official MongoDB Java driver documentation for the latest version.

2. Connect to the MongoDB server:

```
import com.mongodb.client.MongoClient;
import com.mongodb.client.MongoClients;
import com.mongodb.client.MongoDatabase;

public class MongoDBExample {
    public static void main(String[] args) {
        String connectionString = "mongodb://localhost:27017"; // Replace
with your MongoDB connection string
        try (MongoClient mongoClient =
MongoClients.create(connectionString)) {
            System.out.println("Connected to MongoDB");

            // Access a specific database
            MongoDatabase database = mongoClient.getDatabase("mydatabase");
        } catch (Exception e) {
            System.err.println("Error connecting to MongoDB: " +
e.getMessage());
        }
    }
}
```

Replace `mongodb://localhost:27017` with your actual MongoDB connection string, which includes the necessary credentials, database name, and host information.

3. Perform database operations:

Once connected, you can perform various operations like inserting, updating, deleting, and querying documents in MongoDB.

Here's an example of inserting a document into a collection:

```

import org.bson.Document;
import com.mongodb.client.MongoCollection;

public class MongoDBExample {
    public static void main(String[] args) {
        // Connecting to MongoDB...

        // Access a specific collection
        MongoCollection<Document> collection =
database.getCollection("person");

        // Create a document
        Document person = new Document("last_name", "Doe")
            .append("first_name", "John")
            .append("email", "john.doe@example.com")
            .append("age", 25)
            .append("DateOfBr", new Date())
            .append("address", new Document("street", "123 Main
St").append("zip_code", "12345"))
            .append("profession", "Engineer");

        // Insert the document
        collection.insertOne(person);
        System.out.println("Person inserted");
    }
}

```

You can explore other MongoDB operations like updating documents, deleting documents, and querying documents using the MongoDB Java driver's API. Refer to the official MongoDB Java driver documentation for more details: <https://mongodb.github.io/mongo-java-driver/>

4. Disconnect from the MongoDB server:

The MongoDB Java driver manages connections automatically, so you don't need to explicitly disconnect. However, you can do so if necessary:

```

mongoClient.close();

```

By following these steps, you can use MongoDB with a Java application using the MongoDB Java driver.

Spring Boot

To use MongoDB with a Spring Boot application, you can leverage the Spring Data MongoDB module, which provides convenient abstractions and utilities for integrating MongoDB into your application. Here's a step-by-step guide on how to use MongoDB with a Spring Boot application:

1. Set up your Spring Boot project:

Set up a new Spring Boot project or use an existing one. Ensure that you have the necessary dependencies in your pom.xml (Maven) or build.gradle (Gradle) file.

For Maven, add the following dependencies:

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-data-mongodb</artifactId>
</dependency>
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-web</artifactId>
</dependency>
````
```

For Gradle, add the following dependencies:

```
````groovy
implementation 'org.springframework.boot:spring-boot-starter-data-
mongodb'
implementation 'org.springframework.boot:spring-boot-starter-web'
````
```

2. Configure MongoDB connection:

In your application.properties file, specify the MongoDB connection details:

```
spring.data.mongodb.uri=mongodb://localhost:27017/mydatabase
````
```

Replace `'mongodb://localhost:27017/mydatabase'` with your actual MongoDB connection string, including the necessary credentials, database name, and host information.

3. Create a domain model:

Define a Java class representing your MongoDB document. For example, let's create a Person class:

```

import org.springframework.data.annotation.Id;
import org.springframework.data.mongodb.core.mapping.Document;

@Document(collection = "person")
public class Person {
    @Id
    private String id;
    private String lastName;
    private String firstName;
    private String email;
    private int age;
    private Date dateOfBirth;
    private Address address;
    private String profession;

    // Constructors, getters, and setters
}

public class Address {
    private String street;
    private String zipCode;

    // Constructors, getters, and setters
}
...

The `@Document` annotation specifies the MongoDB collection name for this class.

```

4. Create a repository interface:

Create a repository interface that extends `MongoRepository` for your domain model. This interface will provide CRUD operations and other MongoDB-specific queries. For example:

```

import org.springframework.data.mongodb.repository.MongoRepository;

public interface PersonRepository extends MongoRepository<Person,
String> {
    // Additional custom queries can be defined here
}
...

```

5. Perform database operations:

You can now use the `PersonRepository` to perform CRUD operations on the `Person` collection. You can autowire the repository into your service or controller classes to access the MongoDB operations. Here's an example of creating a `Person` document:


```

import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.stereotype.Service;

@Service
public class PersonService {
    private final PersonRepository personRepository;

    @Autowired
    public PersonService(PersonRepository personRepository) {
        this.personRepository = personRepository;
    }

    public void createPerson(Person person) {
        personRepository.save(person);
    }
}
...

```

You can also use other methods provided by `MongoRepository`, such as `findById`, `findAll`, `update`, and `delete`.

6. Build and run the Spring Boot application:

With the above setup, you can now build and run your Spring Boot application, and it will automatically connect to the configured MongoDB instance. You can use the provided CRUD operations to interact with the MongoDB database.

By following these steps, you can use MongoDB with a Spring Boot application using the Spring Data MongoDB module.

PHP

To use MongoDB with PHP, you can utilize the MongoDB extension, which provides a PHP driver for MongoDB. Here's a step-by-step guide on how to use MongoDB with PHP:

1. Install the MongoDB extension:

Make sure you have the MongoDB extension installed and enabled in your PHP environment. You can check the official PHP documentation for installation instructions specific to your operating system:

<https://www.php.net/manual/en/mongodb.installation.php>

2. Connect to the MongoDB server:

```
$manager = new MongoDB\Driver\Manager("mongodb://localhost:27017");

// Access a specific database
$databaseName = "mydatabase";
$database = $manager->selectDatabase($databaseName);
```
```

Replace `'mongodb://localhost:27017'` with your actual MongoDB connection string, which includes the necessary credentials, database name, **and** host information.

### 3. Perform database operations:

Once connected, you can perform various operations like inserting, updating, deleting, and querying documents in MongoDB.

Here's an example of inserting a document into a collection:

```
$collectionName = "person";

// Create a document
$person = [
 "last_name" => "Doe",
 "first_name" => "John",
 "email" => "john.doe@example.com",
 "age" => 25,
 "DateOfBr" => new MongoDB\BSON\UTCDateTime(),
 "address" => [
 "street" => "123 Main St",
 "zip_code" => "12345"
],
 "profession" => "Engineer"
];

// Insert the document
$bulk = new MongoDB\Driver\BulkWrite();
$bulk->insert($person);
$manager->executeBulkWrite("$databaseName.$collectionName", $bulk);
```
```

You can explore other MongoDB operations like updating documents, deleting documents, **and** querying documents using the MongoDB PHP driver's API. Refer to the official MongoDB PHP documentation for more details: <https://docs.mongodb.com/drivers/php/>

4. Close the MongoDB connection:

To close the MongoDB connection, you can simply unset the `$manager` object:

```
unset($manager);
```
```

By following these steps, you can use MongoDB with PHP using the MongoDB extension.

---

## 11- MongoDB Atlas

MongoDB Atlas is a fully managed cloud database service provided by MongoDB, Inc. It offers a convenient and scalable way to deploy, operate, and scale MongoDB databases in the cloud. MongoDB Atlas is designed to simplify the database management process, allowing developers to focus more on building applications and less on database infrastructure.

Key features and aspects of MongoDB Atlas include:

### 1. **Cloud-Based Database Hosting:**

- MongoDB Atlas allows you to host your MongoDB databases in the cloud. It supports major cloud providers such as AWS (Amazon Web Services), Azure (Microsoft Azure), and GCP (Google Cloud Platform).

### 2. **Fully Managed Service:**

- MongoDB Atlas is a fully managed service, meaning that MongoDB, Inc. takes care of various administrative tasks such as database setup, configuration, patching, and backups. This allows developers to focus on application development without the burden of managing the underlying database infrastructure.

### 3. **Automated Scaling:**

- MongoDB Atlas provides automated scaling options, allowing you to easily scale your databases vertically or horizontally as your application's requirements change. You can adjust the instance size, add read-only nodes, or even enable sharding for horizontal scaling.

### 4. **High Availability and Disaster Recovery:**

- MongoDB Atlas offers built-in features for high availability and disaster recovery. It automatically replicates data across multiple nodes and regions, ensuring data durability and availability in the event of hardware failures or other issues.

### 5. **Security Features:**

- MongoDB Atlas includes various security features, such as data encryption in transit and at rest, network isolation, authentication mechanisms, and role-based access control (RBAC). These features help in securing your database and comply with industry standards.

### 6. **Monitoring and Performance Insights:**

- MongoDB Atlas provides monitoring and performance insights through an integrated dashboard. You can view metrics related to database performance, query execution, and resource utilization. Alerts can be set up to notify you of potential issues.

### 7. **Backup and Restore:**

- Automated backups are performed regularly, and you can configure retention policies based on your specific needs. MongoDB Atlas also allows for point-in-time recovery, making it easy to restore your database to a specific moment in time.

## 8. Integration with Other MongoDB Tools:

- MongoDB Atlas seamlessly integrates with other MongoDB tools and services, such as MongoDB Compass (a graphical user interface for MongoDB) and MongoDB Stitch (a serverless platform for building applications).

## 9. Global Clusters:

- MongoDB Atlas supports global clusters, allowing you to deploy databases across multiple geographic regions. This is useful for applications with a global user base, providing low-latency access to data.

## 10. Pay-as-You-Go Pricing:

- MongoDB Atlas follows a pay-as-you-go pricing model. You pay for the resources you use, and pricing is transparent, with no hidden fees. There are different pricing tiers based on the features and resources you need.

Overall, MongoDB Atlas is a robust and user-friendly solution for deploying MongoDB databases in the cloud, offering the benefits of scalability, high availability, security, and ease of management.

---

## How can I use it

Using MongoDB Atlas involves interacting with your database through various means, including the command-line interface (CLI), the MongoDB Compass graphical user interface (GUI), and integrating with tools like Visual Studio Code (VSCode). Here's a brief overview of how you can use MongoDB Atlas with each of these methods:

### 1. Command-Line Interface (CLI):

- Connecting to MongoDB Atlas using MongoDB Shell (mongosh):

1. Install mongosh if you haven't already. You can install it using npm:

```
npm install -g mongosh
```

2. Retrieve your MongoDB Atlas connection string from the Atlas dashboard.
3. Open your terminal and connect to MongoDB Atlas using the connection string:

```
mongosh "<your_connection_string>"
```

Replace <your\_connection\_string> with your actual connection string.

4. Once connected, you can run MongoDB commands, queries, and manage your database using the MongoDB Shell.

### 2. MongoDB Compass:

MongoDB Compass is a graphical user interface for MongoDB that allows you to visually explore your data and manage your databases.

1. Install MongoDB Compass from the official MongoDB website.
2. Open MongoDB Compass and click on "Connect to Host."
3. Enter your MongoDB Atlas connection string in the connection dialog.
4. MongoDB Compass will connect to your Atlas cluster, and you can explore your databases, collections, and run queries through the GUI.

### 3. Visual Studio Code (VSCode):

- Using the MongoDB extension for VSCode:
1. Install the "MongoDB for VSCode" extension from the VSCode marketplace.
  2. Open VSCode and navigate to the "View" menu. Select "Open View..." and choose "MongoDB."
  3. Click on the "Add Connection" button and enter your MongoDB Atlas connection details, including the connection string.
  4. After connecting, you can explore databases, collections, run queries, and interact with your MongoDB Atlas cluster directly from within VSCode.

These are general steps, and the exact details might vary depending on updates to the tools. Always refer to the official documentation for the specific tools you're using for the most accurate and up-to-date information.

Remember to handle connection strings, authentication details, and other sensitive information securely. Always follow best practices for securing your MongoDB Atlas instance, such as using strong passwords, enabling encryption, and setting up appropriate access controls.

---

## 12- MongoDB Express (GUI)

Mongo Express is a web-based administrative interface for MongoDB. It provides a graphical user interface (GUI) that allows users to interact with their MongoDB databases, inspect collections, manage documents, and perform various administrative tasks. Mongo Express is often used as a lightweight and user-friendly alternative to the MongoDB command line interface.

Key features of Mongo Express include:

1. **Web-based Interface:** Mongo Express is accessible through a web browser, providing a graphical and user-friendly interface for managing MongoDB databases.
2. **Database Exploration:** Users can explore the structure of their MongoDB databases, view collections, and examine individual documents within those collections.
3. **Document Editing:** Mongo Express allows users to add, edit, and delete documents within collections. This makes it easier to perform CRUD (Create, Read, Update, Delete) operations without using the MongoDB shell.

4. **Index Management:** Users can manage indexes on collections to optimize query performance.
5. **User Authentication and Authorization:** Mongo Express can be configured to provide authentication and authorization mechanisms, ensuring that only authorized users can access and modify the database.
6. **Query Execution:** It allows users to execute MongoDB queries and view the results directly in the interface.
7. **Server Stats:** Mongo Express provides information about the MongoDB server, including performance statistics, current operations, and server logs.

It's important to note that Mongo Express is a standalone application and needs to be installed separately from MongoDB. It's typically used during development and testing phases and might not be suitable for production environments where more robust security measures are required.

To use Mongo Express, you need to install it and configure it to connect to your MongoDB server. It's important to secure the Mongo Express installation, especially if it's accessible over the internet, to prevent unauthorized access to your MongoDB databases.

## Mongo-Express NodeJs installation

To install Mongo Express, you can follow these general steps. Keep in mind that the specific instructions might vary depending on your operating system and the package manager you use. Below are instructions for installing Mongo Express using npm, a popular package manager for Node.js. Before proceeding, ensure that you have Node.js and npm installed on your system.

1. **Install Node.js and npm:**

If you don't have Node.js and npm installed, download and install them from the official website: [Node.js Downloads \(https://nodejs.org/en/download/\)](https://nodejs.org/en/download/).

2. **Create a new Node.js project:**

Create a new directory for your project and navigate to it in the terminal.

```
mkdir my-mongo-express-app
cd my-mongo-express-app
```

3. **Initialize a new Node.js project:**

Run the following command to create a package.json file for your project.

```
npm init -y
```

4. **Install Mongo Express:**

Use npm to install Mongo Express.

```
npm install mongo-express
```

This installs Mongo Express and its dependencies.

5. **Configure Mongo Express:**

Create a configuration file for Mongo Express. You can either use the sample configuration file provided by Mongo Express or create your own. For simplicity, you can copy the sample configuration.

```
cp node_modules/mongo-express/config.default.js config.js
```

Open `config.js` in a text editor and update the MongoDB connection details to match your setup.

#### 6. **Run Mongo Express:**

Start Mongo Express using the following command:

```
npm start
```

By default, Mongo Express will be accessible at `http://localhost:8081` in your web browser.

#### 7. **Access Mongo Express:**

Open your web browser and navigate to `http://localhost:8081`. You should see the Mongo Express login screen.

Note: By default, there is no authentication. In a production environment, you should secure your Mongo Express installation with authentication and possibly run it behind a reverse proxy with HTTPS.

Remember to consult the official documentation for Mongo Express for any additional or updated installation instructions:

- [Mongo Express GitHub Repository \(https://github.com/mongo-express/mongo-express\)](https://github.com/mongo-express/mongo-express)

Additionally, be cautious when exposing administrative interfaces like Mongo Express to the public internet, and take appropriate security measures to protect your MongoDB database.

---

## 13- Deploying MongoDB with docker

To deploy MongoDB with the latest version using Docker Compose, you can follow these steps:

1. Install Docker: Ensure that Docker and Docker Compose are installed on your machine. Refer to the Docker documentation for instructions specific to your operating system.
2. Create a Docker Compose file: Create a new file named `docker-compose.yml` and open it in a text editor.
3. Define the MongoDB service: In the `docker-compose.yml` file, define the MongoDB service using the following configuration:

```
version: '3'
services:
 mongodb:
 image: mongo
 ports:
 - 27017:27017
 volumes:
 - ./data:/data/db
```

This configuration specifies that you want to use the MongoDB image from Docker Hub, map the container's port 27017 to the host machine's port 27017, and create a volume to persist the MongoDB data in the `./data` directory relative to the `docker-`

compose.yml file.

4. Save the docker-compose.yml file.
5. Start the MongoDB container: Open a terminal or command prompt and navigate to the directory containing the docker-compose.yml file. Run the following command to start the MongoDB container:

```
docker-compose up -d
```

This command will start the containers defined in the docker-compose.yml file in detached mode (-d flag), which means they will run in the background.

6. Verify the MongoDB container: To check if the MongoDB container is running, execute the following command:

```
docker-compose ps
```

You should see the running MongoDB container in the list.

7. Connect to MongoDB: You can connect to the MongoDB instance running inside the container using a MongoDB client or a MongoDB GUI tool. Use the following connection string:

```
mongodb://localhost:27017
```

This assumes that you are running the MongoDB container on the same machine where Docker is installed. If you are running Docker on a different machine, replace "localhost" with the appropriate IP address or hostname.

That's it! You have successfully deployed MongoDB with the latest version using Docker Compose. You can now interact with MongoDB using the connection string and perform database operations as needed. The data will be persisted in the ./data directory on your host machine.

## Using Docker Compose

To deploy MongoDB with the latest version using Docker Compose, you can follow these steps:

1. Install Docker: Ensure that Docker and Docker Compose are installed on your machine. Refer to the Docker documentation for instructions specific to your operating system.
2. Create a Docker Compose file: Create a new file named docker-compose.yml and open it in a text editor.
3. Define the MongoDB service: In the docker-compose.yml file, define the MongoDB service using the following configuration:



```
version: '3'
services:
 mongodb:
 image: mongo
 ports:
 - 27017:27017
 volumes:
 - ./data:/data/db
```

This configuration specifies that you want to use the MongoDB image from Docker Hub, map the container's port 27017 to the host machine's port 27017, and create a volume to persist the MongoDB data in the `./data` directory relative to the `docker-compose.yml` file.

4. Save the `docker-compose.yml` file.
5. Start the MongoDB container: Open a terminal or command prompt and navigate to the directory containing the `docker-compose.yml` file. Run the following command to start the MongoDB container:

```
docker-compose up -d
```

This command will start the containers defined in the `docker-compose.yml` file in detached mode (`-d` flag), which means they will run in the background.

6. Verify the MongoDB container: To check if the MongoDB container is running, execute the following command:

```
docker-compose ps
```

You should see the running MongoDB container in the list.

7. Connect to MongoDB: You can connect to the MongoDB instance running inside the container using a MongoDB client or a MongoDB GUI tool. Use the following connection string:

```
mongodb://localhost:27017
```

This assumes that you are running the MongoDB container on the same machine where Docker is installed. If you are running Docker on a different machine, replace "localhost" with the appropriate IP address or hostname.

That's it! You have successfully deployed MongoDB with the latest version using Docker Compose. You can now interact with MongoDB using the connection string and perform database operations as needed. The data will be persisted in the `./data` directory on your host machine.

## Mongo-express

To deploy MongoDB Express (`mongo-express`) along with MongoDB using Docker Compose, you can follow these steps:

1. Create a Docker Compose file: Create a new file named `docker-compose.yml` and open it in a text editor.

2. Define the services for MongoDB and MongoDB Express: In the `docker-compose.yml` file, define the services for MongoDB and MongoDB Express using the following configuration:

```
version: '3.3'
services:
 mongo:
 image: mongo
 container_name: dkrcomp-mongo
 ports:
 - '27017:27017'
 restart: always
 logging:
 options:
 max-size: 1g
 environment:
 - MONGO_INITDB_ROOT_USERNAME=mongoadmin
 - MONGO_INITDB_ROOT_PASSWORD=bdung

 mongo-express:
 image: mongo-express
 container_name: dkrcomp-mongo-express
 ports:
 - '8081:8081'
 depends_on:
 - mongo
 environment:
 - ME_CONFIG_MONGODB_SERVER=mongo
 - ME_CONFIG_MONGODB_PORT=27017
 - ME_CONFIG_MONGODB_ENABLE_ADMIN=false
 - ME_CONFIG_MONGODB_AUTH_DATABASE=admin
 - ME_CONFIG_MONGODB_AUTH_USERNAME=mongoadmin
 - ME_CONFIG_MONGODB_AUTH_PASSWORD=bdung
 - ME_CONFIG_BASICAUTH_USERNAME=mongoexpress
 - ME_CONFIG_BASICAUTH_PASSWORD=securepassword
```

This configuration specifies two services: `mongodb` and `mongo-express`. The `mongodb` service uses the official MongoDB image, maps the container's port 27017 to the host machine's port 27017, and creates a volume to persist the MongoDB data. The `mongo-express` service uses the official MongoDB Express image, maps the container's port 8081 to the host machine's port 8081, and sets the MongoDB connection details using environment variables.

3. Save the `docker-compose.yml` file.
4. Start the MongoDB and MongoDB Express containers: Open a terminal or command prompt and navigate to the directory containing the `docker-compose.yml` file. Run the following command to start the containers:

```
docker-compose up -d
```

This command will start the containers defined in the `docker-compose.yml` file in detached mode (`-d` flag), which means they will run in the background.

5. Verify the containers: To check if the containers are running, execute the following command:

```
docker-compose ps
```

You should see the running MongoDB and MongoDB Express containers in the list.

6. Access MongoDB Express: Open a web browser and navigate to `http://localhost:8081`. You should see the MongoDB Express interface, where you can manage your MongoDB databases and collections.

That's it! You have successfully deployed MongoDB and MongoDB Express using Docker Compose. You can now access MongoDB Express through the provided URL and perform database management tasks. The MongoDB data will be persisted in the `./data` directory on your host machine.

---